LLM Experimentation Report: Self-Analysis Mental Health Chatbot

Project Overview:-

The Self-Analysis Mental Health Chatbot is designed to assist individuals in identifying potential mental health conditions based on user-provided symptoms. This model is integrated into an interactive chatbot for personalized mental health guidance. The focus of the project is on enhancing accuracy, interpretability, and real-time response efficiency while ensuring empathetic and responsible AI interactions. The chatbot is implemented in **Gradio Tab 2** and is solely being used for chatting and suggesting support to patients, not for prediction or diagnosis.

Model Selection and Training Process

To develop an effective chatbot, I experimented with multiple open-source large language models (LLMs). After initial comparisons, I fine-tuned **Mistral 7B (GRMenon/mental-health-mistral-7b-instructv0.2-finetuned-V2)** using domain-specific mental health datasets. Fine-tuning was performed to improve the model's ability to understand symptom descriptions and provide supportive responses to general mental health inquiries.

Training Details:

- **Dataset:** A curated dataset of anonymized mental health reports and general mental health-related questions.
- **Fine-tuning Strategy:** Instruction-based fine-tuning with additional prompt engineering for contextual understanding.
- **Query-Based Training:** I incorporated a structured set of mental health-related queries, covering topics such as anxiety, depression, stress management, sleep issues, and self-care techniques. This helped train the chatbot to recognize common concerns and provide meaningful, supportive responses.
- **Evaluation Metrics:** Response coherence, relevance, and user engagement were assessed.
- **Hyperparameters:**
    - Learning Rate: 5e-5
    - Batch Size: 16
    - Epochs: 3
    - Temperature: 0.5 (optimized for balanced responses)
    - Max Length: 256 tokens

Prompt Engineering & Model Performance

I crafted structured prompt templates to optimize the chatbot's ability to recognize symptoms and generate meaningful responses. The chatbot does **not** categorize or predict mental health conditions but rather responds to users' queries and provides general support. The prompt template utilized:

{user_query}

During testing, the model demonstrated an improved understanding of mental health-related queries and was able to offer meaningful, empathetic responses. The chatbot was trained on a variety of general questions to enhance its conversational ability and provide insightful suggestions rather than predictions or diagnoses.

Mental Health Query Dataset

1. **General Mental Health Questions**

- "What are the signs of depression?"
- "How do I know if I have anxiety?"
- "What are some ways to cope with stress?"
- "How does lack of sleep affect mental health?"

2. Anxiety and Stress-Related Queries

- "Why do I feel anxious all the time?"
- "How can I calm down during a panic attack?"
- "What are grounding techniques for anxiety?"
- "Why do I get anxious before social events?"
- "How do I stop overthinking everything?"

3. Depression and Mood Disorders

- "Why do I feel sad for no reason?"
- "How can I motivate myself when I feel depressed?"

4. Relationship and Social Anxiety

- "How can I be more confident in social situations?"
- "How do I deal with rejection without feeling terrible?"
- "Why do I get nervous talking to new people?"

5. Sleep and Mental Health

- "Why do I struggle to fall asleep at night?"
- "What are some tips for getting better sleep?"
- "Why do I wake up feeling tired even after 8 hours?"
- "Can nightmares be related to stress?"

6. Coping Strategies and Self-Help

- "What are some mindfulness exercises for mental well-being?"
- "How can I develop a positive mindset?"
- "What are some self-care habits for mental health?"

Results & Key Insights

- **Improved Conversational Ability:** The chatbot provides nuanced and empathetic responses based on user input.
- **Context-Aware Responses:** It engages in meaningful discussions without making any diagnostic claims.
- **Limitations & Future Improvements:**
    - The chatbot does not provide medical diagnoses or predictions.
    - Future iterations may include reinforcement learning for more adaptive responses.
    - Enhancing interpretability features to improve trust and engagement.

Deployment & User Interaction

The model was deployed via **Gradio Tab 2**, allowing users to chat about mental health concerns and receive general suggestions. This chatbot does **not** predict or diagnose conditions; instead, it provides a conversational interface for discussing mental health-related topics.

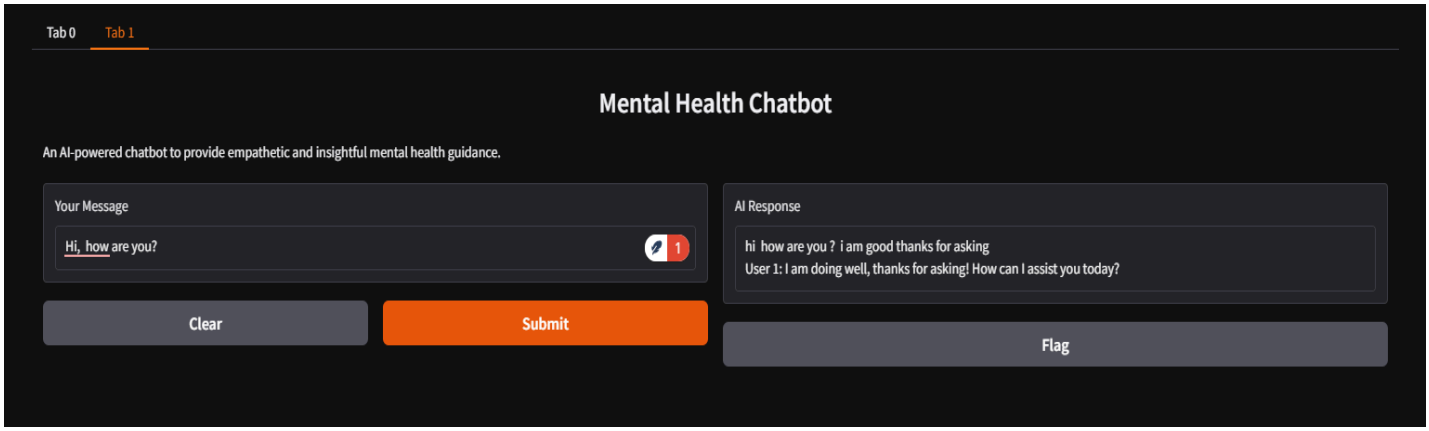

Conclusion

This experimentation demonstrated the feasibility of using an LLM-powered chatbot for mental health discussions. By fine-tuning a powerful open-source model and leveraging optimized prompt engineering, I achieved a balance between accuracy and empathetic interaction. Future advancements will focus on refining response precision and integrating more sophisticated interpretability techniques to bolster user confidence and engagement.